

Machine Learning Engineer Nanodegree

Image Classification with Transfer Learning by implementing InceptionV3 and MobileNetV2

Jeffen Chen
September 8th, 2018

I. Definition

Project Overview

Neural networks have revolutionized many areas, such as natural language processing, image classification, and autonomous driving. However, as neural networks become more powerful, their complexity and training efficiency also pose challenges for humans. How to balance between accuracy and training efficiency of neural networks is a hot topic in recent years. Many excellent models and algorithms have also been developed, such as AlexNet, VGGNet, GoogLeNet, Inception. These models and algorithms can achieve good results but take a lot of time and effort to train.

In practice, very few people train an entire Convolutional Network from scratch (with random initialization), because it is relatively rare to have a dataset of sufficient size. Instead, it is common to pretrain a ConvNet on a very large dataset (e.g. ImageNet, which contains 1.2 million images with 1000 categories), and then use the ConvNet either as an initialization or a fixed feature extractor for the task of interest.

Problem Statement

In the project, we will use the transfer learning method to train an existing large neural network for image classification. We expect the new model to have higher accuracy and training efficiency than a simple neural network built manually. Keras already comes with some deep neural network models and trained weights. We will choose one model for pre-training and fine tuning. We will redefine the input layer and the fully connected layer associated with the output to better fit the image size and output we enter. The ultimate goal of this project is to use the trained model to classify the CIFAR-100.

Metrics

The evaluation metric for this competition is multi-class classification accuracy i.e. the proportion of true class labels correctly predicted. Because the distribution of each type of image in the project is uniform, so accuracy can be calculated under following equation:

$$accuracy = \frac{True\ Predictions}{All\ Predictions} \times 100\%$$

II. Analysis

Data Exploration

I decided to choose CIFAR-100 dataset after research. This dataset is publicly available at University of Toronto website and it has images in 100 classes containing 600 images each. There are 500 training images and 100 testing images per class. The 100 classes in the CIFAR-100 are grouped into 20 superclasses. The 100 object class labels are shown in fig.1. 50000 labelled examples are provided for training, with a further 10000 unlabelled examples used for testing. Each images has 3 RGB colour channels and pixel dimensions 32×32 for an

Algorithms and Techniques

Transfer learning is introduced in this project. I used InceptionV3 with weights pre-trained on ImageNet, and MobileNetV2 initialized randomly which built manually according to the paper by Mark Sandler.

The top layer of InceptionV3 is replaced with the following structure:

global_average_pooling2d_1 (Glo	(None, 2048)	0	mixed10[0][0]
dense_1 (Dense)	(None, 1024)	2098176	global_average_pooling2d_1[0][0]
dropout_1 (Dropout)	(None, 1024)	0	dense_1[0][0]
dense_2 (Dense)	(None, 100)	102500	dropout_1[0][0]
=====			

Figure 3: InceptionV3 top layer structure

In the pre-training of InceptionV3, we freeze all layers except the top layer. In fine tuning, I want to keep the low-level features and only train the advanced features, so I only freeze the first two blocks of InceptionV3 and reduce the learning rate.

In MobileNetV2, I just added one fully-connected layer as the top layer to output 100 categories.

The complete MobileNetV2 structure is shown in figure 4. Each line describes a sequence of 1 or

Input	Operator	t	c	n	s
$224^2 \times 3$	conv2d	-	32	1	2
$112^2 \times 32$	bottleneck	1	16	1	1
$112^2 \times 16$	bottleneck	6	24	2	2
$56^2 \times 24$	bottleneck	6	32	3	2
$28^2 \times 32$	bottleneck	6	64	4	2
$28^2 \times 64$	bottleneck	6	96	3	1
$14^2 \times 96$	bottleneck	6	160	3	2
$7^2 \times 160$	bottleneck	6	320	1	1
$7^2 \times 320$	conv2d 1x1	-	1280	1	1
$7^2 \times 1280$	avgpool 7x7	-	-	1	-
$1 \times 1 \times k$	conv2d 1x1	-	k	-	-

Figure 4: MobileNetV2 structure

more identical (modulo stride) layers,

repeated n times. All layers in the same

sequence have the same number c of

output channels. The first layer of each

sequence has a stride s and all others use

stride 1. All spatial convolutions use

3×3 kernels. The expansion factor t is

always applied to the input size.

Benchmark

The benchmark model will be a simple CNN consist 3 fully-connected layers and a dense layer that shown in the figure.2. This will be relatively simple model to train and probably will get very low score. But it will still be more accurate than most supervised machine learning algorithms. Using this benchmark, you can more clearly compare the advantages of transfer learning in terms of operational efficiency and results relative to CNN.

Layer (type)	Output Shape	Param #
conv2d_98 (Conv2D)	(None, 224, 224, 16)	208
max_pooling2d_8 (MaxPooling2	(None, 112, 112, 16)	0
conv2d_99 (Conv2D)	(None, 112, 112, 32)	2080
max_pooling2d_9 (MaxPooling2	(None, 56, 56, 32)	0
conv2d_100 (Conv2D)	(None, 56, 56, 64)	8256
max_pooling2d_10 (MaxPooling	(None, 28, 28, 64)	0
global_average_pooling2d_3 ((None, 64)	0
dense_4 (Dense)	(None, 100)	6500
Total params: 17,044		
Trainable params: 17,044		
Non-trainable params: 0		

Figure 5: Simple CNN model structure

III. Methodology

Data Preprocessing

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape $(nb_samples, rows, columns, channels)$ where $nb_samples$ corresponds to the total number of images (or samples), and rows, columns, and channels correspond to the number of rows, columns, and channels for each image, respectively.

```

from keras.preprocessing import image
from tqdm import tqdm

def path_to_tensor(img_path):
    # loads RGB image as PIL.Image.Image type
    img = image.load_img(img_path, target_size=(224, 224))
    # convert PIL.Image.Image type to 3D tensor with shape (224, 224, 3)
    x = image.img_to_array(img)
    # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and return 4D tensor
    return np.expand_dims(x, axis=0)

def paths_to_tensor(img_paths):
    list_of_tensors = [path_to_tensor(img_path) for img_path in tqdm(img_paths)]
    return np.vstack(list_of_tensors)

```

Figure 6: Code that convert image to tensor

The `path_to_tensor` function above takes a string-valued file path to a colour image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape $(nb_sample, 224, 224, 3)$.

```

def generate(batch, train_tensors, train_targets, valid_tensors, valid_targets):
    # Using the data Augmentation in training data
    datagen1 = ImageDataGenerator(
        shear_range=0.2,
        rotation_range=90,
        width_shift_range=0.2,
        height_shift_range=0.2,
        horizontal_flip=True)

    datagen2 = ImageDataGenerator(rotation_range=90)

    train_generator = datagen1.flow(
        train_tensors,
        train_targets,
        batch_size=batch)

    validation_generator = datagen2.flow(
        valid_tensors,
        valid_targets,
        batch_size=batch)

    count1 = len(train_tensors)
    count2 = len(valid_tensors)
    return train_generator, validation_generator, count1, count2

```

Figure 7: Image augmentation configuration

After the image is converted to tensor, we can use ImageDataGenerator for image augmentation. In this project, we rotate, scale, cut and flip the image, and use image flipping for the validation data. The image before and after conversion is shown in the figure below. However, in practice, since image augmentation will greatly reduce the training efficiency, this method is not used in real training process.

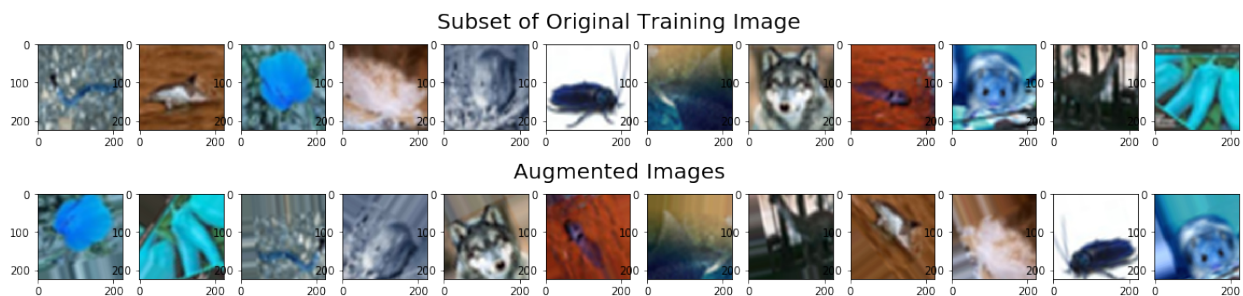


Figure 8: Original training images and augmented images

Implementation

InceptionV3

Because the Keras already includes InceptionV3 and the trained ImageNet weights, all we need to do is introduce the model and replace fully-connected layer. In pre-training, all layers except the fully-connected layer need to be frozen. In fine tuning, we only need to freeze the first 249 layers (2 blocks). Also, the optimizer and learning rate used in both trainings are different. In fine tuning we want to avoid overfitting with a lower learning rate. The code looks like figure 9.

MobileNetV2

The implementation of the model is mainly based on given structure in the paper which also been shown in figure 4. For coding detail see figure 10. Same as InceptionV3, earlyStopping and checkpoint is used to save time and prevent model from over-fitting.

```

def Inception(input_shape, k):
    # create the base pre-trained model
    base_model = InceptionV3(input_shape=input_shape,
                             weights='imagenet', include_top=False)

    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dense(1024, activation='relu')(x)
    x = Dropout(0.3)(x)
    # and a logistic layer -- let's say we have 100 classes(k)
    predictions = Dense(k, activation='softmax')(x)

    model = Model(inputs=base_model.input, outputs=predictions)

    for layer in base_model.layers:
        layer.trainable = False

    return model

def fine_tune(model):
    model.load_weights('saved_models/weights.best.inception.hdf5')

    for layer in model.layers[:249]:
        layer.trainable = False
    for layer in model.layers[249:]:
        layer.trainable = True

    return model

def train(batch, epochs, num_classes, size, weights, train_tensors, train_targets, valid_tensors, valid_targets):
    if weights:
        model = Inception((size, size, 3), num_classes)
        model = fine_tune(model)
        model.compile(optimizer=SGD(lr=0.0001, momentum=0.9), loss='categorical_crossentropy')
    else:
        model = Inception((size, size, 3), num_classes)
        model.compile(optimizer='rmsprop', loss='categorical_crossentropy')

    earlystop = EarlyStopping(monitor='val_loss', patience=15, verbose=1, mode='auto')
    checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.inception.hdf5',
                                   verbose=1, save_best_only=True)

    model.fit(train_tensors, train_targets,
              validation_data=(valid_tensors, valid_targets),
              epochs=epochs, verbose=2, batch_size=batch, callbacks=[earlystop, checkpointer])

```

Figure 8: InceptionV3 implementation


```

def _conv_block(inputs, filters, kernel, strides):
    """Convolution Block
    This function defines a 2D convolution operation with BN and relu6.
    # Arguments
        inputs: Tensor, input tensor of conv layer.
        filters: Integer, the dimensionality of the output space.
        kernel: An integer or tuple/list of 2 integers, specifying the
            width and height of the 2D convolution window.
        strides: An integer or tuple/list of 2 integers,
            specifying the strides of the convolution along the width and height.
            Can be a single integer to specify the same value for
            all spatial dimensions.
    """

    channel_axis = 1 if K.image_data_format() == 'channels_first' else -1

    x = Conv2D(filters, kernel, padding='same', strides=strides)(inputs)
    x = BatchNormalization(axis=channel_axis)(x)
    return Activation(relu6)(x)

def _bottleneck(inputs, filters, kernel, t, s, r=False):
    """Bottleneck
    # Arguments
        inputs: Tensor, input tensor of conv layer.
        filters: Integer, the dimensionality of the output space.
        kernel: An integer or tuple/list of 2 integers, specifying the
            width and height of the 2D convolution window.
        t: Integer, expansion factor.
            t is always applied to the input size.
        s: An integer or tuple/list of 2 integers, specifying the strides
            of the convolution along the width and height. Can be a single
            integer to specify the same value for all spatial dimensions.
        r: Boolean, Whether to use the residuals.
    """

    channel_axis = 1 if K.image_data_format() == 'channels_first' else -1
    tchannel = K.int_shape(inputs)[channel_axis] * t

    x = _conv_block(inputs, tchannel, (1, 1), (1, 1))

    x = DepthwiseConv2D(kernel, strides=(s, s), depth_multiplier=1, padding='same')(x)
    x = BatchNormalization(axis=channel_axis)(x)
    x = Activation(relu6)(x)

    x = Conv2D(filters, (1, 1), strides=(1, 1), padding='same')(x)
    x = BatchNormalization(axis=channel_axis)(x)

    if r:
        x = add([x, inputs])
    return x

def _inverted_residual_block(inputs, filters, kernel, t, strides, n):

    x = _bottleneck(inputs, filters, kernel, t, strides)

    for i in range(1, n):
        x = _bottleneck(x, filters, kernel, t, 1, True)

    return x

def MobileNetV2(input_shape, k):

    inputs = Input(shape=input_shape)
    x = _conv_block(inputs, 32, (3, 3), strides=(2, 2))

    x = _inverted_residual_block(x, 16, (3, 3), t=1, strides=1, n=1)
    x = _inverted_residual_block(x, 24, (3, 3), t=6, strides=2, n=2)
    x = _inverted_residual_block(x, 32, (3, 3), t=6, strides=2, n=3)
    x = _inverted_residual_block(x, 64, (3, 3), t=6, strides=2, n=4)
    x = _inverted_residual_block(x, 96, (3, 3), t=6, strides=1, n=3)
    x = _inverted_residual_block(x, 160, (3, 3), t=6, strides=2, n=3)
    x = _inverted_residual_block(x, 320, (3, 3), t=6, strides=1, n=1)

    x = _conv_block(x, 1280, (1, 1), strides=(1, 1))
    x = GlobalAveragePooling2D()(x)
    x = Reshape((1, 1, 1280))(x)
    x = Dropout(0.3, name='Dropout')(x)
    x = Conv2D(k, (1, 1), padding='same')(x)

    x = Activation('softmax', name='softmax')(x)
    output = Reshape((k,))(x)

    model = Model(inputs, output)
    plot_model(model, to_file='images/MobileNetV2.png', show_shapes=True)

    return model

```

Figure 10: MobileNetV2 implementation

Refinement

Tuning the learning rate turned out to have the most significant impact on the result of the project. By adjusting learning rate from 0.01 to 1e-4 in SGD optimizer in InceptionV3 in fine-tuning stage, the accuracy was raised from 46% to nearly 50%. Similar improvement was achieved by reduced learning rate in Adam optimizer which used in MobileNetV2. Also, adding dropout layer to the top layer can improve accuracy in some extend, but dropout value did not seems to have significant impact on the final result.

IV. Results

Model Evaluation and Validation

Training InceptionV3 and MobileNetV2 took me about 8 hours each on AWS p2.xlarge instance. For efficiency, I have to limit training dataset to half of the whole dataset and train model with a relatively medium rate, otherwise it will cost too much time to train. I trained all these models several times with different optimizer and learning rate. The final accuracy on CIFAR-100 test is shown below, the accuracy is based on the best result I got in the training.

	Benchmark CNN	InceptionV3	MobileNetV2
Accuracy	15.3%	49.54%	50.49%
Time per epoch	40s	217s	360s

Justification

From the training results, we can see that the two models finally achieved an accuracy of about 50%, which is 4-5 times better than of benchmark CNN, and the difference is obvious. It is worth

noting that all three models use the normal size of the original data set for training, so if you use more training data and use image preprocessing, we should be able to achieve better results.

After MobileNetV2 reaches the training bottleneck, the loss value will soon start to be discrete, so it is necessary to further reduce the learning rate. But because of time, I did not make further optimizations to the model.

Since the image was exported and resized, some images are distorted (not sure whether it is the cause of cv2), so there should be errors when processing the actual real world image.

V. Conclusion

Free-Form Visualization

Figure 11 shows the confusion matrix of InceptionV3 and MobileNetV2. It would be too small to see clear in this report, you can check out the file in `saved_models/` directory. From the matrix we can see that certain class has better performance compare to others. For example, class 36(mice) and 58(mountain). Also, some class tends to be more easy be recognized as another

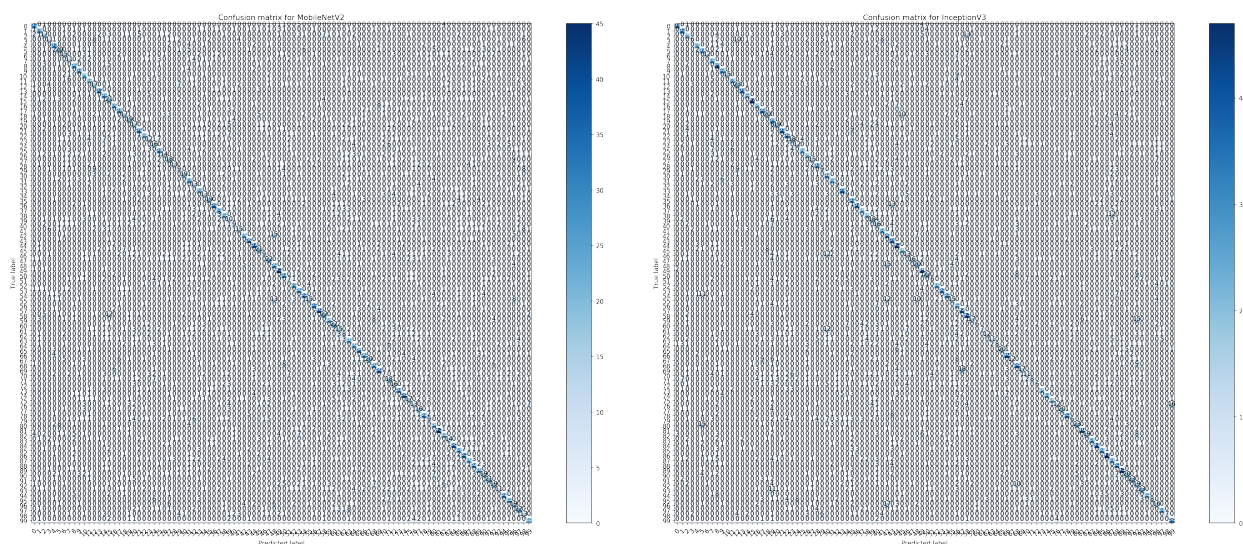


Figure 11: Confusion matrix for two models.

class. For example, class 42(leopard) and 48, class 55(shrew) and 48. I guess the reason behind might because the two classes share the same pixel feature. After we know which classes are easy to cause confusion, we can train our model purposely based on the data we have.

Reflection

I think this is a very challenging project. From the initial selection of research directions, the selection of data sets, the cleaning of data, the development of models and training and debugging all require a lot of research. Because I am not only a beginner in machine learning, but also a beginner in Python, there are many basic knowledge in this process that I need to learn. But I think there are two most difficult points in the whole project. One is read the research papers and build MobileNetV2 by hand, and the other is to choose the right optimizer and parameters to train. These two processes cost me a lot of time to read and debug the model. Due to the computational power limitation, I can't use whole datasets for training, but still achieved a 50% accuracy based on 0.5 datasets. The training efficiency of these two models surprised me because I was able to reach the training bottleneck in a short time. I am looking forward to using this model in other image recognition applications.

Improvement

Due to computational constraints, I only used 50% of the dataset for training and did not apply image augmentation. So using more training data and using image augmentation is an improvement. In addition, whether using InceptionV3 or MobileNetV2, the training quickly reached the bottleneck, val_loss and val_acc no longer improved, and the model gradually overfitting. Choosing and using more appropriate optimizers and parameters, reducing the learning rate should be able to alleviate this situation to some extent. According to the actual training

experience, adjusting the top layer is not effective in improving the accuracy, so this method still needs to continue to explore.

References

- [1] Szegedy, Christian, et al. "Rethinking the inception architecture for computer vision. " *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016.
- [2] Sandler, Mark, et al. "Inverted residuals and linear bottlenecks: Mobile networks for classification, detection and segmentation." arXiv preprint arXiv:1801.04381 (2018).